

Date 14th March 2024

## Quiz-2

Thursday

### → Shared Memory Programming :- Lecture 8 :-

- 1.) physically: processors in a computer share access to the same RAM
- 2.) virtual: threads running on the processors interact with one another thru shared vars in the common addr space of a single proc

⇒ performance improvements to serial code is easier with multi-threading than with message-passing.


→ MP reqs code to be redesigned

→ MT allows incremental parallelism

- Clusters are made up of "multiple processors per compute node"
- Processors interact & synchronize with each other thru shared vars

OpenMP with MPI improves performance

→ Fork/Join :- at 2 levels: shared & distributed



→ MT prog is most common shared-mem programming method

serial code ⇒ the proc is master & begins exe  
or only executing S ⇒ || portion, master S can fork more SSSS  
⇒ || portion finish, SSSS join & Master S continues executing serial code

### Shared-Memory

- one active thread at start & end of prog
- Num of active S inside prog changes dynamically  
process of converting a seq prog to a parallel prog a bit at a time
- Supports incremental parallelism

### Message-Passing

- all processes remain active thru out exe of program
- Seq to Parallel transformation requires major effort  
transformation done in one giant step rather than many tiny steps
- No incremental parallelism



openMP: standard method for shared-memory prog  
MPI: standard for distributed memory prog  
→ Codes are portable  
Date → Performance is good

Openmp consists of compiler directives, lib func, env vars

pragma: pragmatic information

```
#pragma omp parallel for  
for(i=0; i<N; i++)
```

① Compiler should know total num of iterations b4 executing the prog

```
a[i] = b[i] + c[i]
```

③ Loops with 'continue' statement are allowed

② Body of for-loop must not allow premature exits (e.g. break etc)  
× return, exit, goto not allowed

① Shared Variable: same addr in execution context of every thread

② Private Variable: diff addr in execution context of every thread

∴ thread can NOT access priv variables of another thread

∴ C prog express data-parallel operations as for loops

## Lecture 9:-

### Clauses:-

∴ Compiler takes care of generating code that fork/join threads & allocates iterations to threads

- 1.) Private: directs compiler to make one or more variables private  

```
#pragma omp parallel for private(j)
```
- 2.) First Private: create private variables having initial values identical to the value of the variable controlled by the Master Thread as the loop is entered.
- 3.) Last Private: used to copy back to the master thread's copy of the var, the priv copy of the variable from the thread that executed the last iteration.
- 4.) Reduction: takes care of storing partial results in private variables & combining partial results after the loop reduction (+: variable)
- 5.) if-clause: conditional statement, decides at run-time whether the loop should be executed in parallel or not



Date \_\_\_\_\_

6.) Scheduling : scheduling the loops means dividing num of iterations b/w the processes.

- a chunk is a contiguous range of iterations

- increasing chunk size reduces scheduling overhead & may increase cache hit rate (due to operations on contiguous mem loc)

- decreasing chunk size allows finer balancing of workloads

7.) Static Scheduling :

- 1.) if chunk-size given : splits iteration space into equal chunks & assigns them to threads in a round robin fashion

- 2.) if no chunk-size : iteration space is split into as many chunks as there are threads

$$\frac{\text{num of iterations}}{\text{total threads}} : \text{ } \Rightarrow \text{ one chunk is assigned to each thread}$$

- decision about work division is done before code execution

- lower scheduling overhead but can cause load-imbalance if all processors are not of same compute-capability

8.) Dynamic Scheduling :

- 1.) iteration space is partitioned into chunks given by chunk-size

- 2.) every thread is assigned single chunk

- 3.) decision for remaining iteration chunks is done on run-time.

- 4.) Chunk is assigned to threads as they become idle (takes care of temporal imbalances resulting from static scheduling)

⇒ When you use dynamic sched without specifying chunk size, OpenMP assigns one iterations to each thread initially. As threads complete assigned iterations, they request & are assigned new iterations. No single thread remains idle while others have work to do.



Date 30/3/24

mid-2

Sat

## Lecture-9:- (continued)

### → Guided Scheduling :-

- Scheduling algo dynamically adjusts the size of the chunks based on the num of remaining iterations & the execution progress.
- At the beginning → larger chunks of iterations.
- As the parallel loop progresses, threads complete their assigned chunks, guided scheduling gradually decreases the size of the chunks.
- This adaptive beh helps in achieving better load balancing among threads esp as the workload becomes more uneven or as some iterations take longer to execute than others.
- Guided Scheduling allows you to specify a minimum chunk size ('C'). Once the chunk size decreases to this min value, it remains constant for subsequent iterations.  
(this helps in preventing chunk size from becoming too small, which could lead to increased overhead)

1.) exploits initial parallelism by assigning bigger chunks in start

2.) decreased chunk-size exploits fine-grained load balancing


### ⇒ No Wait Clause :

- avoids implicit barrier (every thread waits after executing for loop, waits for the rest of threads)
- thread can easily move to next after completing assigned tasks/ iterations



Date

⇒ functional/task parallelism :-



• `#pragma omp sections [clause list]`

→ if code is based on diff segments or sections

→ can execute a, b, d parallelly

`#pragma omp parallel sections` (creates a team of threads which execute the sections in the region parallelly)

`#pragma omp section`

`v = a()`

Omp parallel sections generates it's own team of threads

`#pragma omp section`

`w = b()`

Omp sections uses existing team of threads

`#pragma omp section`

`x = d()`

3 distributes section among the threads

if multiple sections pragmas are inside one parallel block, may reduce fork/join costs

Basic Lib Functions :-

- Controlling Num of Threads & Processes :-

- 1.) `omp_get_num_procs`  
num of CPUs/cores in machine
- 2.) `omp_get_num_threads`  
num of active threads, call from 11 region
- 3.) `omp_get_max_threads`  
value of env variable `omp_num_threads`  
can be called outside of parallel region
- 4.) `omp_get_thread_num`  
value of thread id
- 5.) `omp_set_num_threads`  
`omp_set_dynamic(0)` disable dynamic  
`omp_set_num_threads(4)`

- 1.) `omp_set_num_threads`
- 2.) `omp_get_num_threads`
- 3.) `omp_get_max_threads`
- 4.) `omp_get_thread_num`
- 5.) `omp_get_num_procs`
- 6.) `omp_in_parallel`

- Controlling & Monitoring { Creation :-

- 1.) `omp_set_dynamic`
- 2.) `omp_get_dynamic`
- 3.) `omp_set_nested`
- 4.) `omp_get_nested`



Date

## Lecture 10 :-

### => Synchronization in OpenMp:-

- 1.) Barrier directive :- all threads in a team wait until others have caught up, & then release (opposite of No Wait & is already implemented default)
- 2.) Single Directive :- a structured block executed by single thread in parallel region (mostly print statements) **implicit barrier**
- 3.) Master Directive :- specialization of single: only master thread will execute structured block **no implicit barrier**
- 4.) Atomic Directive :- specifies that the single memory location update should be performed as an atomic operation (single operation w/o interruption from other threads)
- 5.) Critical Section :- code segment that has a shared var & needs to be executed as an atomic action. Only one process can execute critical section at a time

### => Environment Variables

- |                                  |                               |
|----------------------------------|-------------------------------|
| 1.) <code>omp_num_threads</code> | 3.) <code>omp_schedule</code> |
| 2.) <code>omp_dynamic</code>     | 4.) <code>omp_nested</code>   |

#### 1.) `Omp_Dynamic` :-